

Exp.No: 2

Run a basic Word Count Map Reduce program to understand Map Reduce Paradigm

AIM:

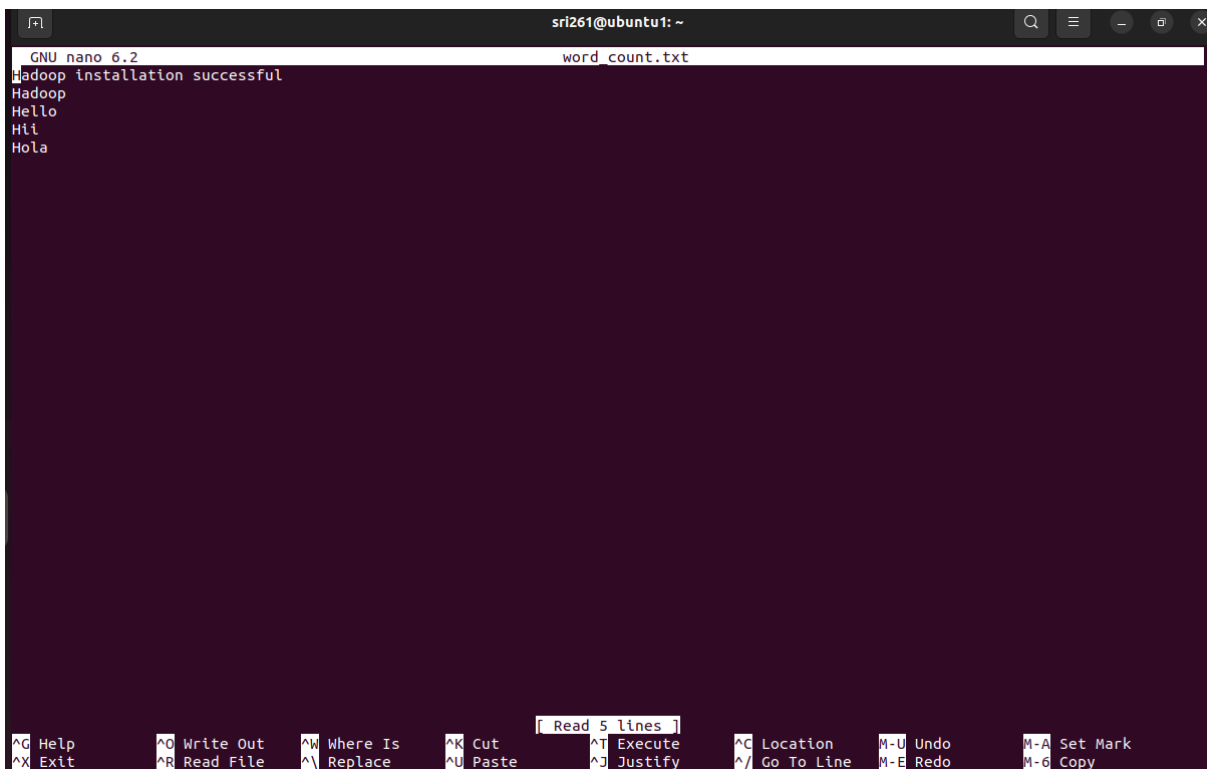
To run a basic Word Count MapReduce program.

Procedure:**Step 1: Create Data File:**

Create a file named "word_count_data.txt" and populate it with text data that you wish to analyse. Login with your hadoop user.

```
nano word_count.txt
```

Output: Type the below content in word_count.txt



```
sri261@ubuntu1: ~
GNU nano 6.2 word_count.txt
Hadoop installation successful
Hadoop
Hello
Hii
Hola

^G Help      ^O Write Out  ^W Where Is   ^K Cut        ^T Execute    ^C Location   M-U Undo      M-A Set Mark
^X Exit      ^R Read File  ^\ Replace    ^U Paste      ^J Justify    ^_ Go To Line M-E Redo      M-6 Copy
```

Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

```
nano mapper.py
# Copy and paste the mapper.py code

#!/usr/bin/env python3
# import sys because we need to read and write data to STDIN and STDOUT
#!/usr/bin/python3
import sys
for line in sys.stdin:
    line = line.strip() # remove leading and trailing whitespace
    words = line.split() # split the line into words
    for word in words:
        print( '%s\t%s' % (word, 1))
.
```

Step 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

```
nano reducer.py
# Copy and paste the reducer.py code
```

reducer.py

```
#!/usr/bin/python3 from operator
import itemgetter import sys
current_word = None current_count
= 0 word = None for line in
sys.stdin: line = line.strip()
word, count = line.split("\t", 1)
try:
    count = int(count)
except ValueError:
    continue
    if current_word
== word: current_count
+= count else:
    if current_word:
        print( '%s\t%s' % (current_word, current_count))
    current_count = count current_word = word if
current_word == word: print( '%s\t%s' %
(current_word, current_count))
```

Step 4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data.

```
start-all.sh hdfsdfs -mkdir /word_count_in_python hdfsdfs -copyFromLocal
/path/to/word_count.txt/word_count_in_python
```

Step 6: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

```
chmod 777 mapper.py reducer.py
```

Step 7: Run Word Count using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the Word Count program using Hadoop Streaming.

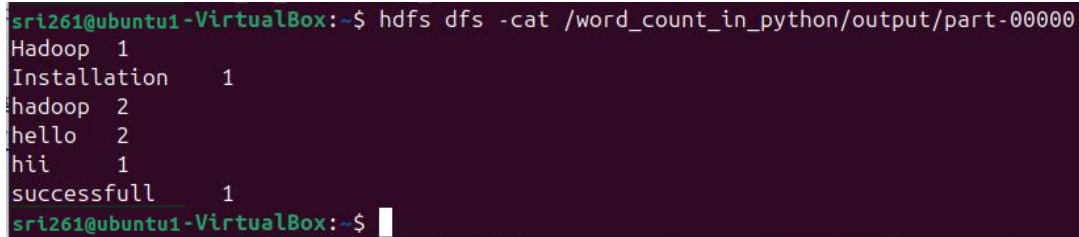
```
hadoop jar /path/to/hadoop-streaming-3.3.6.jar \ -input
/path/to/word_count_in_python/word_count_data.txt \
-output /word_count_in_python/new_output \
-mapper /path/to/mapper.py \
-reducer /path/to/reducer.py
```

```
sri261@ubuntu1-VirtualBox:~$ hadoop jar /home/hadoop/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.3.6.
jar -input /word_count_in_python/word_count.txt -output /word_count_in_python/output -mapper mapper.py -r
educer reducer.py
2024-09-02 10:50:26,297 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-09-02 10:50:26,431 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-09-02 10:50:26,431 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2024-09-02 10:50:26,440 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2024-09-02 10:50:26,716 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-02 10:50:26,813 INFO mapreduce.JobSubmitter: number of splits:1
2024-09-02 10:50:26,983 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1711765797_0001
2024-09-02 10:50:26,984 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-09-02 10:50:27,133 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2024-09-02 10:50:27,134 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2024-09-02 10:50:27,135 INFO mapreduce.Job: Running job: job_local1711765797_0001
2024-09-02 10:50:27,160 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutput
tCommitter
2024-09-02 10:50:27,188 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2024-09-02 10:50:27,188 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary fold
ers under output directory:false, ignore cleanup failures: false
2024-09-02 10:50:27,262 INFO mapred.LocalJobRunner: Waiting for map tasks
2024-09-02 10:50:27,264 INFO mapred.LocalJobRunner: Starting task: attempt_local1711765797_0001_m_000000_
0
2024-09-02 10:50:27,316 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2024-09-02 10:50:27,320 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary fold
ers under output directory:false, ignore cleanup failures: false
2024-09-02 10:50:27,357 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2024-09-02 10:50:27,368 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/word_count_in_python
/word_count.txt:0+62
2024-09-02 10:50:27,412 INFO mapred.MapTask: numReduceTasks: 1
2024-09-02 10:50:27,450 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
```

Step 8: Check Output:

Check the output of the Word Count program in the specified HDFS output directory.

```
hdfs dfs -cat /word_count_in_python/new_output/part-00000
```

A terminal window with a dark purple background. The prompt is 'sri261@ubuntu1-VirtualBox:~\$'. The command 'hdfs dfs -cat /word_count_in_python/output/part-00000' has been entered. The output is: 'Hadoop 1', 'Installation 1', 'hadoop 2', 'hello 2', 'hii 1', 'successfull 1'. The prompt is now 'sri261@ubuntu1-VirtualBox:~\$' with a cursor.

```
sri261@ubuntu1-VirtualBox:~$ hdfs dfs -cat /word_count_in_python/output/part-00000
Hadoop 1
Installation 1
hadoop 2
hello 2
hii 1
successfull 1
sri261@ubuntu1-VirtualBox:~$
```

Result:

Thus, the program for basic Word Count Map Reduce has been executed successfully.